Roadmap — phases end at proof gates

Extracted from Section 3.6 of the final report. Phase 2 does not start because Phase 1 finished; it starts because Phase 1 met a measured threshold.

Phase	Focus	Proof gate before proceeding
0 — done	Deterministic platform: submission spine, fees, tasks, approvals, campaigns, exams, audit. Gemini narrative summaries and meeting intelligence live.	Platform adopted as the daily system of record; daily submission compliance stable.
1 — MVP pilot (1 month, 5 campuses)	Daily operations triage in assist-only mode with the full evaluation harness.	Severity recall, summary error rate, override rate and reviewer trust thresholds all met.
2 — controlled expansion (1 quarter)	All 40 campuses; cross-campus weekly pattern detection; OCR document intake; budget-proposal completeness checking; scoped natural-language query.	Cost per campus per month within the modelled ceiling; no campus-scoping breach; sustained reviewer agreement rate.
3 — assisted exam operations	Paper equivalence assistance and portion-completion inference feeding the examinations flow, with mandatory faculty approval.	Zero AI-only paper approvals; faculty acceptance measured; academic governance sign-off obtained.
4 — evaluate productisation	Assess whether the platform is offered to peer institutions as a product.	Two consecutive quarters of stable internal operation and a defensible support model.

Cut rather than sequenced

Two high-value ideas were removed from the plan entirely rather than deferred to a later phase, because cost of error — not value — decided them:

- Predictive fee-default and dropout risk scoring. The training data available today reflects operational recording habits more than student behaviour, and a wrong prediction attaches to a named student.
- Automated staff performance scoring from operational logs. It converts an operations tool into a surveillance instrument, which would destroy the trust the platform depends on for honest daily reporting.

Prerequisites for Phase 2, independent of model quality

- Documentation, a written runbook and a second engineer. One builder is not an operating model, and this is the most serious operational risk in the submission.
- Data minimisation in place before any text reaches the model API.
- A working single-flag rollback to rules-only ordering.